


Optimization Project - Giovanni Pelosi

Heating and Electricity Consumption Prediction

Artificial Neural Network
Approach

Heating and Electricity Consumption Prediction



❖ Artificial Neural Network


- ❖ Introduction
- ❖ Training
- ❖ Stochastic Gradient Descent
- ❖ ADAM
- ❖ Deep Neural Networks (RNN and LSTM)

❖ Project


- ❖ Data structure
- ❖ Theoretical goal
- ❖ Experiments
- ❖ Results
- ❖ Considerations and Conclusions

❖ References

Artificial Neural Network



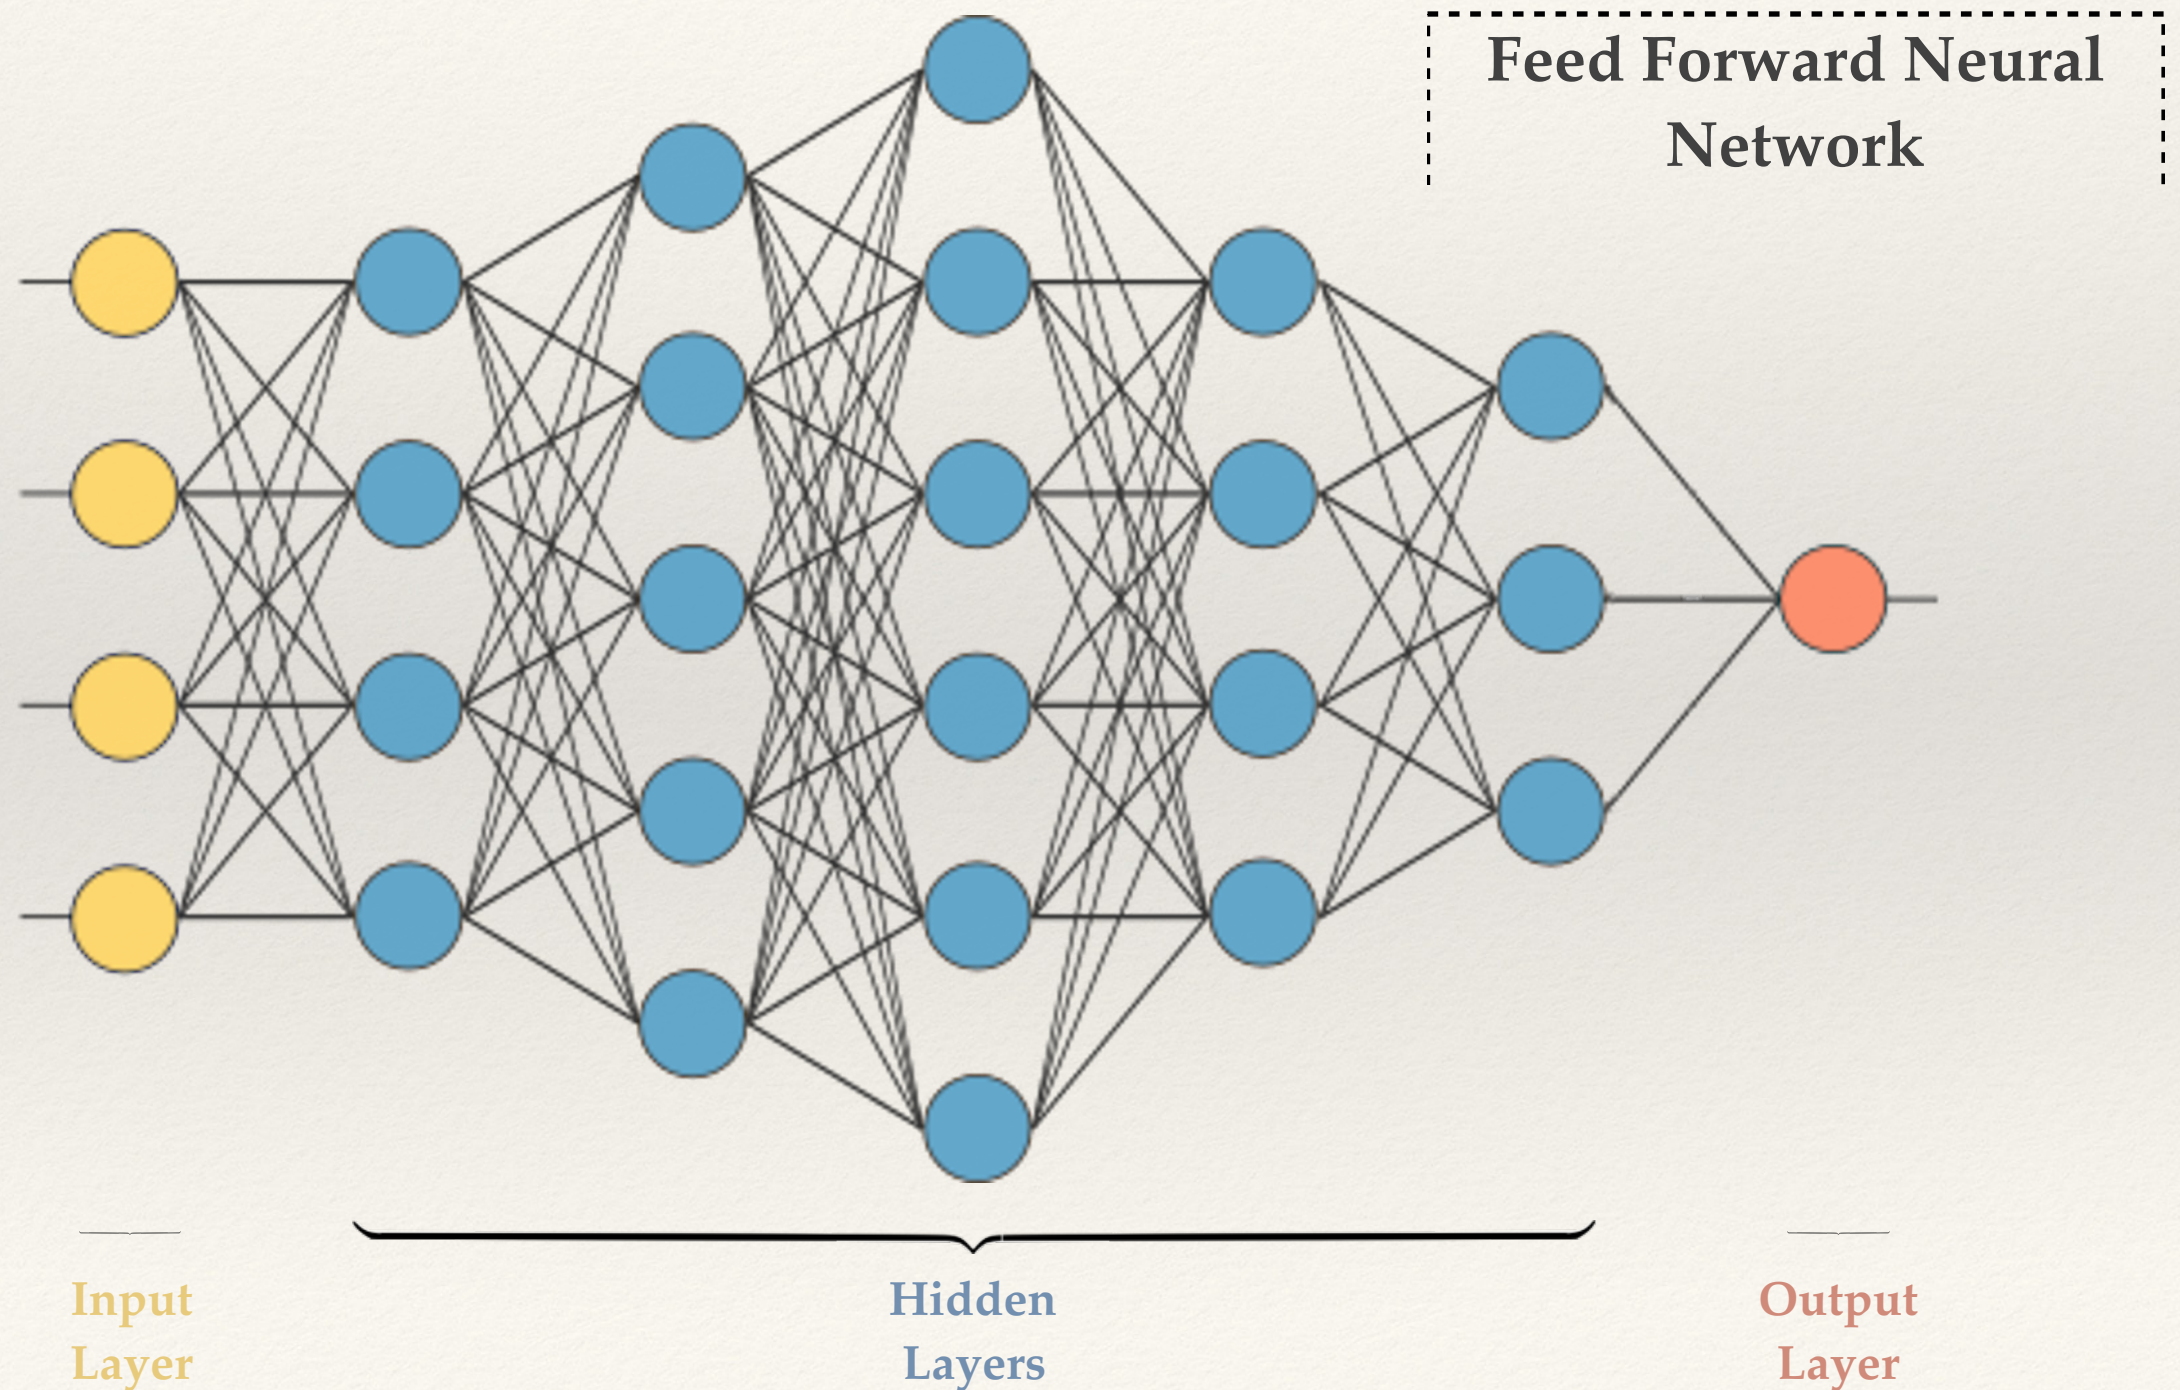
An ANN is a net made up by **Artificial Neurones**




- Weights w_{ji}
- Activation value z_j
- Activation threshold (bias) b_j
- Activation function $g_j(\cdot)$ [Sign, Linear, Sigmoid, Hyperbolic tangent, ReLU, ...]

Artificial Neural Network


An ANN is a net made up by **Artificial Neurones**



Training an ANN



- ❖ Back-propagation (Gradient Descent) $w^{new} := w^{old} - \eta \frac{\partial E}{\partial w}$
- ❖ 2 phase algorithm (1 propagation, 2 weight update)



$$y = g\left(\sum_j W_j \cdot h\left(\sum_i w_{ji} \cdot x_i\right)\right)$$


$$E = \sum_n (t_n - y_n)^2$$

$$a_j = \sum_i^I w_{ji} \cdot x_i \quad b_j = h(a_j) \quad A = \sum_j^J W_j b_j$$

$$\begin{aligned} \frac{\partial E}{\partial W_j} &= \sum_n^N 2(t - g(A)) \cdot \frac{\partial}{\partial W_j} (t - g(A)) \\ &= \sum_n^N 2(t - g(A)) \cdot (-g'(A)) \cdot \frac{\partial}{\partial W_j} A \\ &= \sum_n^N 2(t - g(A)) \cdot (-g'(A)) \cdot b_j \end{aligned}$$

$$W_j^{k+1} = W_j^k + 2\eta \sum_n^N (t - g(A)) \cdot g'(A) \cdot b_j$$

Issues in Training ANN



❖ Convergence

- Gradient Descent with Momentum
- Quasi Newton Methods
- Conjugate Gradient


❖ Local Optima

- Multiple Restart
- Randomised Algorithms

❖ Generalisation and Overfitting


- Early Stopping
- Weight Decay

SGD Algorithm




- ❖ Online Algorithm: Each input generates a weight adjustment
- ❖ Mini-Batches in practice
- ❖ If the learning rate decreases with an appropriate rate, SGD converges to a global minimum when the objective function is convex (to a local minimum otherwise)

Improving SGD



- ❖ **Momentum:** remembers the update Δw at each iteration and determines the next update as a linear combination of the gradient and the previous update
- ❖ **AdaGrad** *Adaptive Gradient Algorithm:* SGD with per-parameter learning rate [increases the learning rate for more sparse parameters and decreases the learning rate for less sparse ones]
- ❖ **RMSProp** *Root Mean Square Propagation:* learning rate is adapted for each of the parameters. It divides the learning rate of a weight by a running average of the magnitudes of recent gradients for that weight.


ADAM algorithm



Adaptive Moment Estimation


- ❖ Combination of AdaGrad and RMSProp
- ❖ Invariant to diagonal rescaling of the gradients
- ❖ Suited for problems that are large in terms of data and / or parameters
- ❖ Adapts the parameter learning rates based on the average first moment (the mean) as in RMSProp, and on the average of the second moments of the gradients (the uncentered variance)

Deep Neural Networks




- ❖ FeedForward Neural Networks FFNNs
- ❖ Recurrent Neural Networks RNNs
- ❖ Long Short Term Memory Networks LSTM
- ❖ Convolutional Neural Networks CNNs

Recurrent Neural Networks

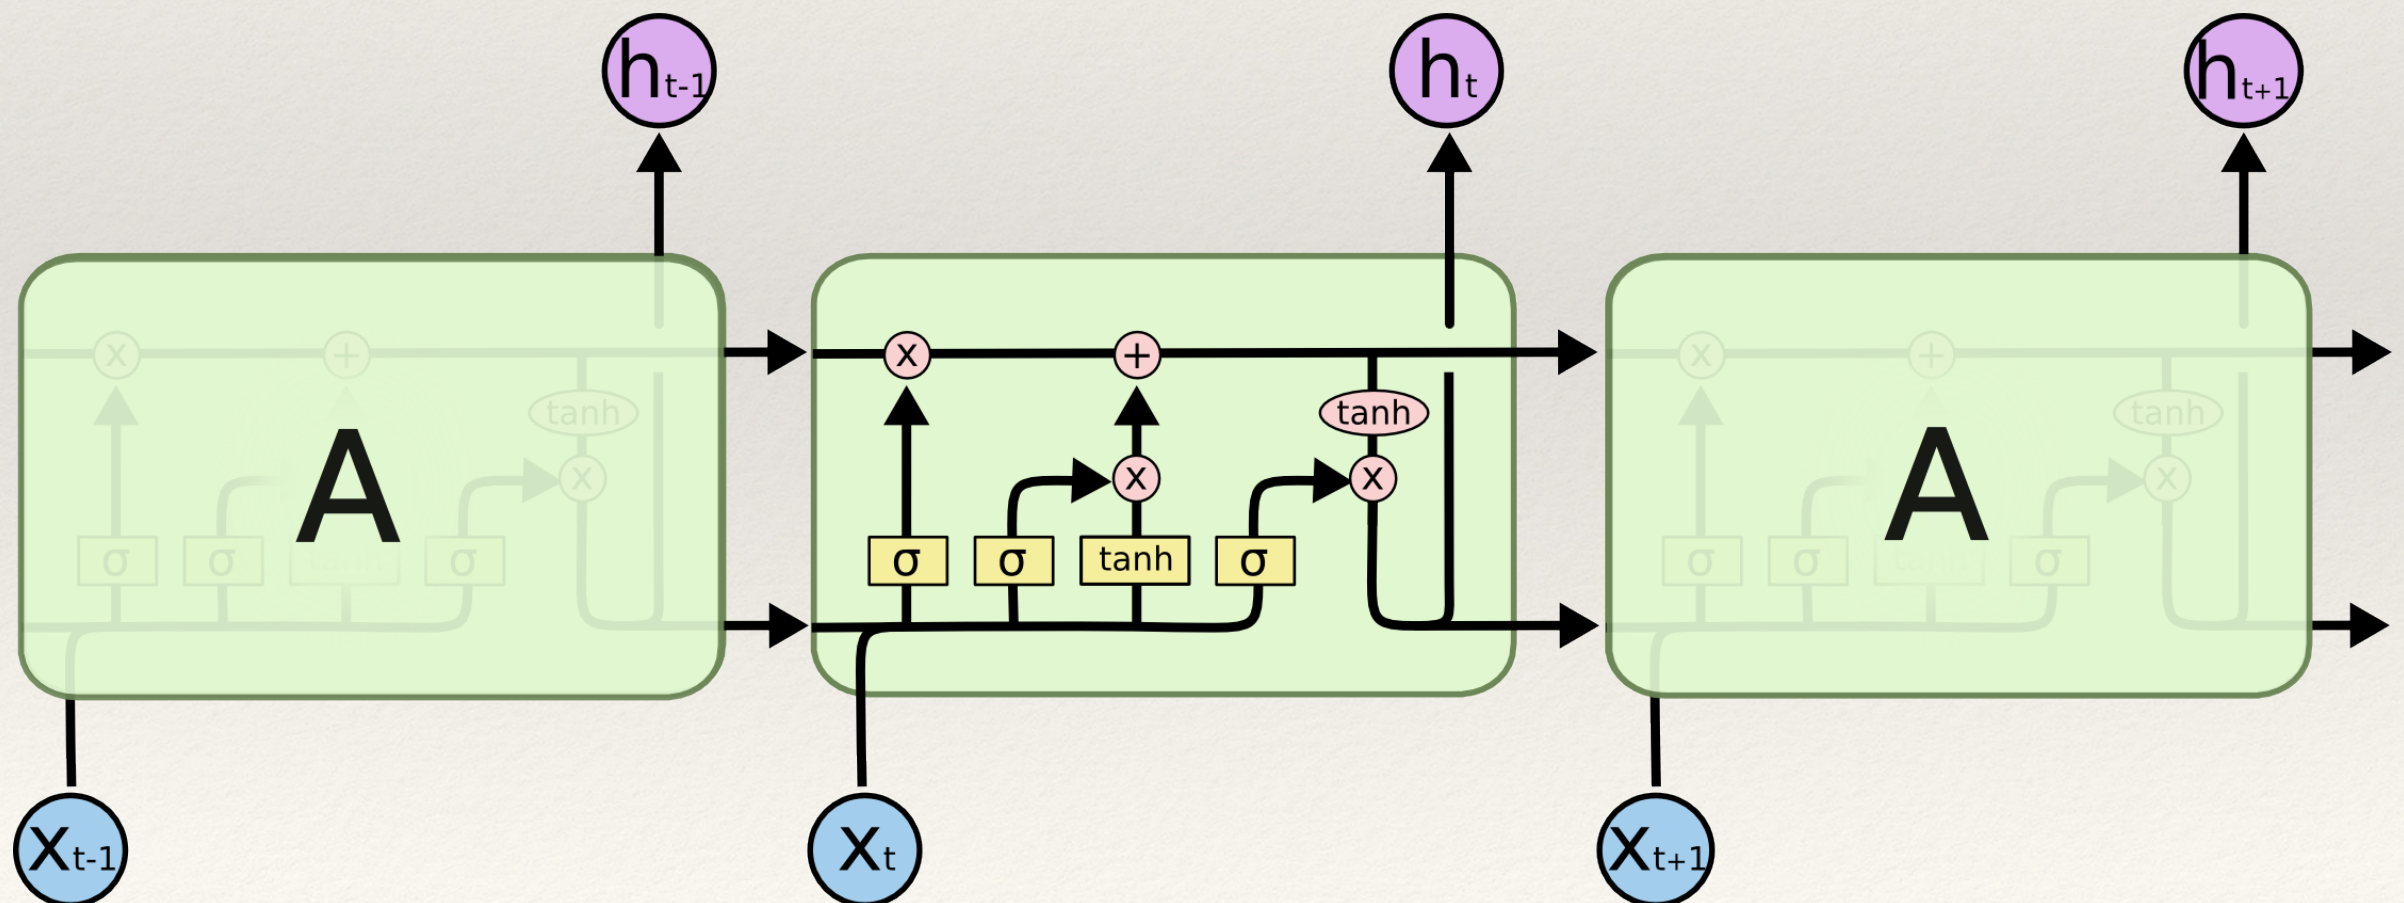


- ❖ Add a new unit b to the hidden layer and a new input unit $c(t)$ to represent the value of b at time $(t - 1)$. b thus can summarise information from earlier values of x arbitrarily distant in time



Long Short Term Memory



- ❖ Special kind of RNN, capable of learning long-term dependencies





Long Short Term Memory



Cell State



Long Short Term Memory



$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

Forget Gate

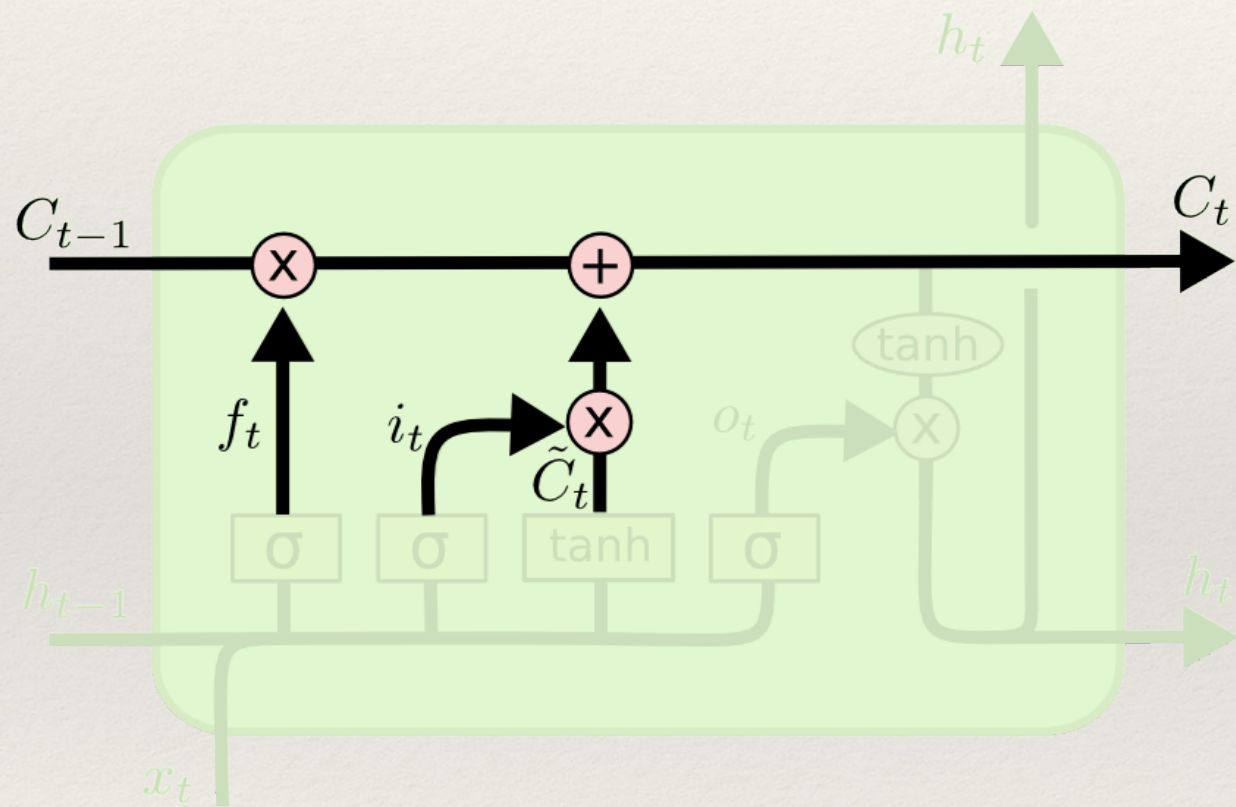

Long Short Term Memory



$$i_t = \sigma (W_i \cdot [h_{t-1}, x_t] + b_i)$$
$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Input Gate

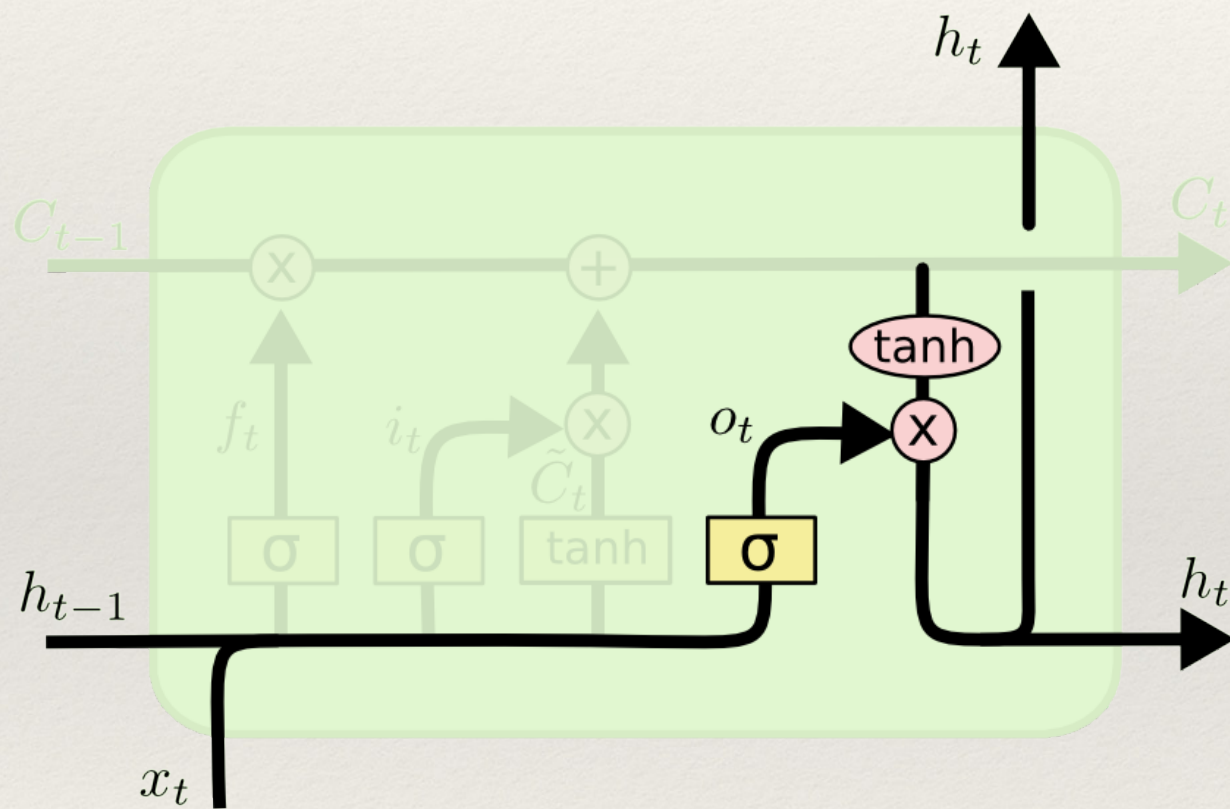

Long Short Term Memory



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

Memory Gate

Long Short Term Memory




$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh (C_t)$$

Output Gate


PROs and CONs of DNNs



Pros	Cons
Important features are learned automatically from data	A lot of data is needed
No domain knowledge required	Hard to analyse - Black box
	Initialization


Project Goal

Heating and Electricity Consumption Prediction

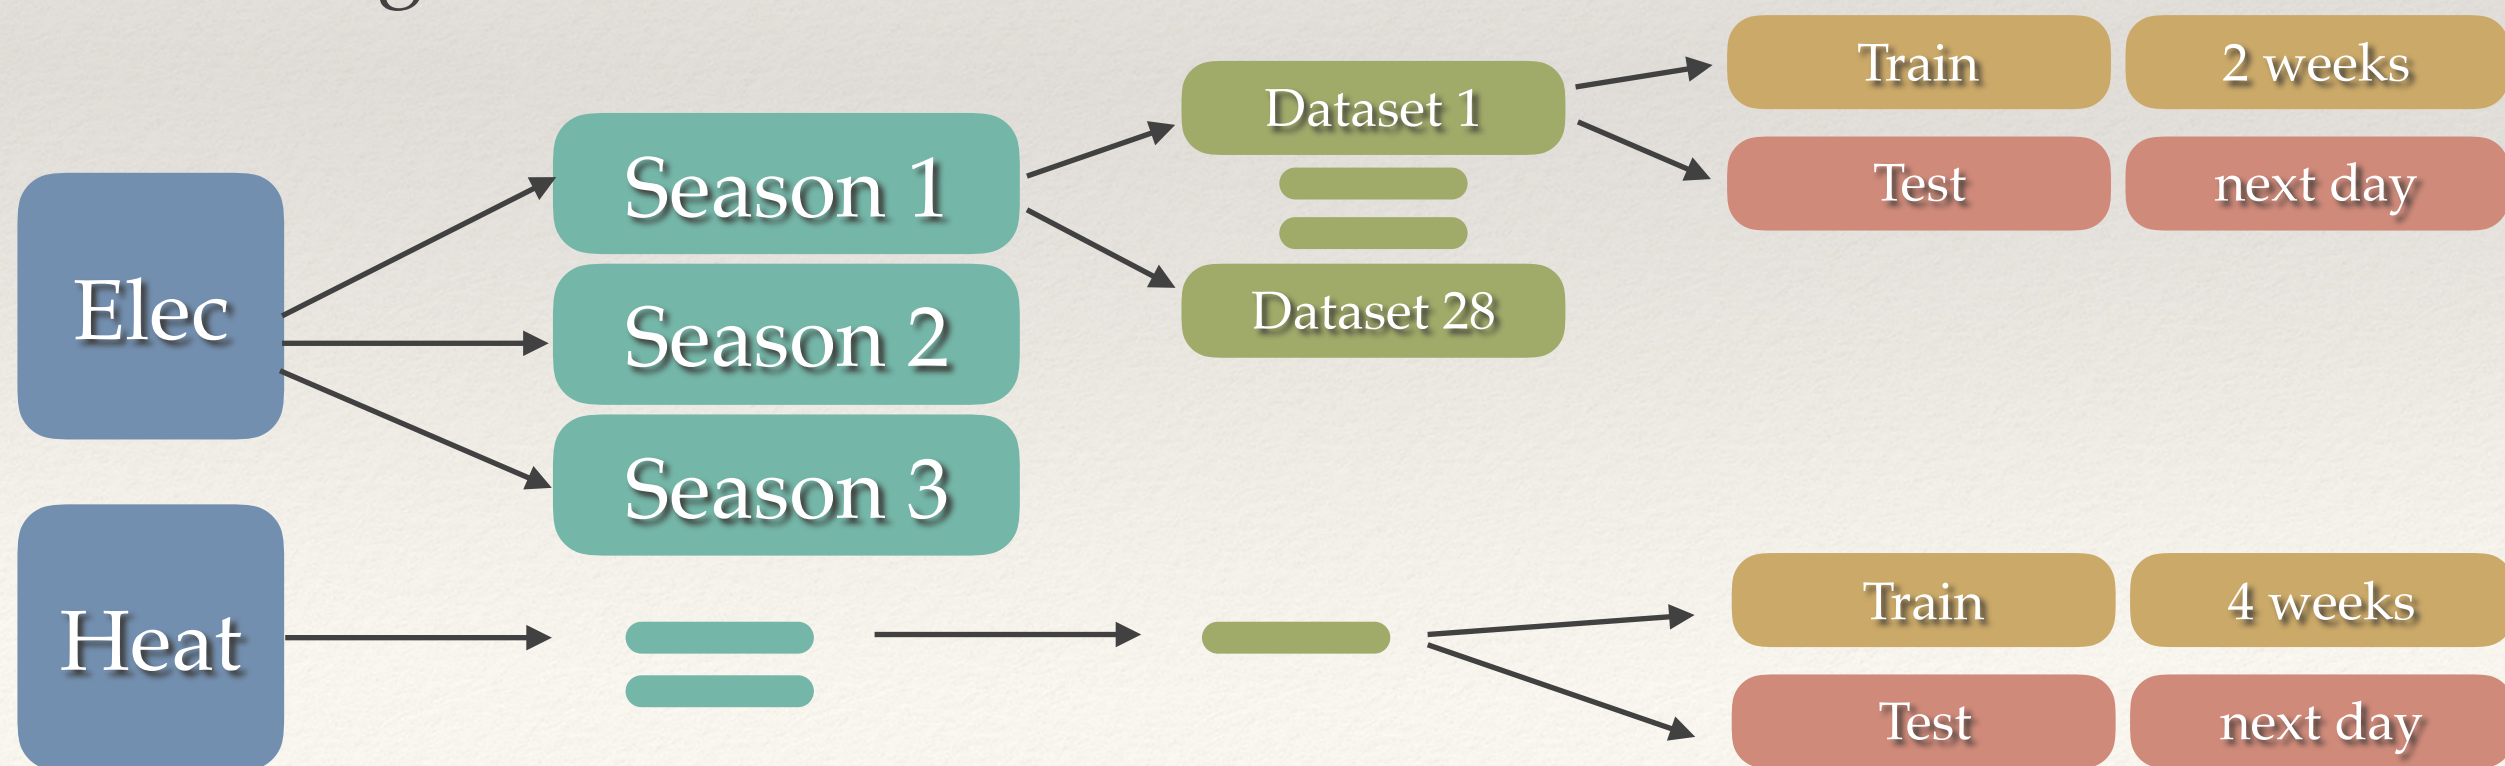


- ❖ Predict the next 24 hours of Heating and Electricity consumption of an hospital.


Data (1)



- ❖ 2 datasets : one for the heating data and one for the electricity data
- ❖ Each dataset is divided in 3 sub-datasets (3 seasons)
- ❖ Each season is composed by 28 small dataset used for training and testing




Data (2)



❖ 11 features + 15 features


Original	Additional
4 weather forecast infos for the specific hour (sample)	12 weather forecast infos of the 3 previous (hours) samples
6 binary variables that identify the week day	3 variables indicating the electric (or heating) consumption of the same hour of the same week day of the previous three weeks
1 binary variable that states if the day is holiday or not	

Theoretical Goal




- ❖ To discover if a DNN trained on the "11 features" dataset is able to predict with a accuracy similar to that of a single hidden layer NN trained on handcrafted features (11 + 15). Is the DNN able to learn the important features that characterise the given problem?
- ❖ Compare the results with a LSTM NN trained on the 11 features. Is a LSTM NN able to exploit the past to better predict the future?

Experiments



1. First experiment with simple FFNN
2. Preprocess the dataset by removing seasonal components
3. Modify the structure and the training options of the nets
4. *~~Data multiplication~~*
5. Enlarge the dataset (adding 2 weeks in the training [slide_horizon and incremental])

Best Results



ELECTRICITY Top 10 - 2 weeks dataset


	Season 1		Season 2		Season 3	
	kind of net :	rmse	kind of net :	rmse	kind of net :	rmse
1	"all_deseason_40"	"0.059021"	"small_deseason_70"	"0.11623"	"all_deseason_56"	"0.055899"
2	"all_deseason_72"	"0.059078"	"small_deseason_58"	"0.11646"	"all_deseason_80"	"0.055916"
3	"all_deseason_32"	"0.059283"	"small_deseason_82"	"0.11652"	"all_deseason_44"	"0.056047"
4	"all_deseason_44"	"0.059527"	"small_deseason_62"	"0.11675"	"all_deseason_40"	"0.056232"
5	"all_deseason_76"	"0.060081"	"small_deseason_66"	"0.11722"	"all_deseason_36"	"0.056552"
6	"all_deseason_68"	"0.060191"	"small_deseason_74"	"0.11753"	"all_deseason_76"	"0.056567"
7	"all_deseason_16"	"0.060398"	"small_deseason_78"	"0.11814"	"all_deseason_68"	"0.056705"
8	"all_deseason_80"	"0.060421"	"small_57"	"0.1183"	"all_deseason_60"	"0.057028"
9	"all_deseason_8"	"0.061021"	"small_61"	"0.12021"	"all_deseason_64"	"0.057055"
10	"all_deseason_36"	"0.061359"	"small_81"	"0.12337"	"all_deseason_72"	"0.057151"

HEATING Top 10 - 4 weeks dataset

	Season 1		Season 2		Season 3	
	kind of net	rmse	kind of net :	rmse	kind of net :	rmse
1	"all_55"	"0.1058"	"all_75"	"0.071589"	"all_deseason_84"	"0.059433"
2	"all_35"	"0.10647"	"all_43"	"0.07177"	"all_deseason_64"	"0.059747"
3	"all_71"	"0.10706"	"all_55"	"0.071785"	"all_deseason_60"	"0.060165"
4	"all_31"	"0.10754"	"all_3"	"0.072347"	"all_deseason_28"	"0.06032"
5	"all_59"	"0.10768"	"all_47"	"0.072429"	"all_deseason_72"	"0.06074"
6	"all_43"	"0.10825"	"all_35"	"0.072557"	"all_deseason_68"	"0.060958"
7	"all_63"	"0.10903"	"all_31"	"0.072701"	"all_deseason_80"	"0.06116"
8	"all_67"	"0.11012"	"all_67"	"0.073447"	"all_deseason_56"	"0.061252"
9	"all_39"	"0.1102"	"all_79"	"0.073576"	"all_deseason_36"	"0.061338"
10	"all_83"	"0.11024"	"all_51"	"0.073652"	"all_deseason_76"	"0.061376"

From 1 to 56 -> SGD Algorithm ; From 57 to 84 -> ADAM Algorithm

Best Results



ELECTRICITY - 4 weeks Increased dataset


	Season 1		Season 2		Season 3	
	Kind of net	rmse	Kind of net	rmse	Kind of net	rmse
<i>2 weeks Best</i>	"all_deseason_40"	"0.059021"	"small_deseason_70"	"0.11623"	"all_deseason_56"	"0.055899"
1	FFNN_SH_3	"0.07161"	FFNN_I_8	"0.17826"	FFNN_SH_3	"0.054087"
2	FFNN_SH_4	"0.072802"	FFNN_SH_3	"0.18829"	FFNN_SH_2	"0.054209"
3	FFNN_SH_2	"0.073216"	FFNN_SH_2	"0.19272"	FFNN_SH_5	"0.056093"
4	FFNN_SH_5	"0.074144"	FFNN_SH_4	"0.19675"	FFNN_SH_4	"0.059436"
5	FFNN_I_8	"0.075395"	FFNN_SH_1	"0.19687"	FFNN_I_8	"0.061119"
6	FFNN_SH_1	"0.076003"	FFNN_SH_5	"0.20191"	FFNN_SH_1	"0.070659"
7	LSTM_I_6	"0.089365"	LSTM_I_6	"0.22048"	LSTM_I_6	"0.085532"
8	LSTM_I_7	"0.090354"	LSTM_I_7	"0.22324"	LSTM_I_7	"0.10065"

HEATING - 6 weeks Increased dataset

	Season 1		Season 2		Season 3	
	Kind of net	rmse	Kind of net	rmse	Kind of net	rmse
<i>4 weeks Best</i>	"all_55"	"0.1058"	"all_75"	"0.071589"	"all_deseason_84"	"0.059433"
1	FFNN_SH_1	"0.10814"	FFNN_SH_1	"0.087646"	FFNN_SH_1	"0.069542"
2	FFNN_I_8	"0.11132"	FFNN_SH_2	"0.088027"	FFNN_SH_4	"0.071043"
3	FFNN_SH_4	"0.11171"	FFNN_SH_3	"0.091076"	FFNN_SH_5	"0.073204"
4	FFNN_SH_5	"0.11209"	FFNN_SH_4	"0.091672"	FFNN_I_8	"0.075838"
5	FFNN_SH_3	"0.11351"	FFNN_SH_5	"0.091917"	FFNN_SH_2	"0.077845"
6	FFNN_SH_2	"0.11388"	FFNN_I_8	"0.099661"	FFNN_SH_3	"0.078878"
7	LSTM_I_7	"0.34203"	LSTM_I_6	"0.13229"	LSTM_I_7	"0.20192"
8	LSTM_I_6	"0.36638"	LSTM_I_7	"0.14767"	LSTM_I_6	"0.27224"


- SH : Sliding Horizon dataset
- I : incremental dataset
- RMSE : average of the daily error per season

Considerations (1)




- ❖ 26 features vs 11 features / de-trend vs trend / preprocessing

Considerations (2)




❖ Enlarged dataset vs original dataset




Heating 4
weeks
sample

Increased presence of
annual trends




Heating 6
weeks
sample

Considerations (3)



- ❖ LSTM initial problems (slow start) → slow learning



FFNN_SH_1

FFNN_SH_2

FFNN_SH_3

FFNN_SH_4


FFNN_SH_5

LSTM_I_6

LSTM_I_7


FFNN_I_8

Conclusions



- ❖ DNN not always good → need a large amount of “good” training data
- ❖ DNN might not be able to extract the useful features
- ❖ Training execution time vs accuracy
- ❖ Black box model → not easy to find the best structure or training options

References



- ◉ <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>
- ◉ <https://it.mathworks.com/help/nnet/examples/time-series-forecasting-using-deep-learning.html>
- ◉ <https://it.mathworks.com/help/nnet/ug/improve-neural-network-generalization-and-avoid-overfitting.html>
- ◉ <https://it.mathworks.com/help/nnet/ug/ist-of-deep-learning-layers.html>
- ◉ <https://machinelearningmastery.com/improve-deep-learning-performance/>
- ◉ <https://it.mathworks.com/videos/developing-forecast-models-from-time-series-data-in-matlab-part-1-93067.html>
- ◉ <https://it.mathworks.com/videos/developing-forecast-models-from-time-series-data-in-matlab-part-2-93066.html>
- ◉ Diederik P. Kingma, Jimmy Lei Ba, ADAM: A METHOD FOR STOCHASTIC OPTIMIZATION
- ◉ Grippo, Manno, Sciandrone, Decomposition Techniques for Multilayer Perceptron Training
- ◉ LeCun, Bengio, Hinton, Deep learning
- ◉ Glorot, Bengio, Understanding the difficulty of training deep feedforward neural networks
- ◉ Bottou, Large-Scale Machine Learning with Stochastic Gradient Descent